

IT2164/IT2561

Operating Systems and Administration

Chapter 2

Operating System Organization

Operating System Organization

At the end of this chapter, you should be able to

- List and describe the common services offered by an OS
- Describe the functions of an operating system
- Understand the operating system requirements
- Describe the basic operating system organization
- Explain the general implementation issues and implementation mechanisms
- Understand modern OS architecture

Operating System Services

- An operating system provides an environment for the execution of programs
- Different operating systems are organized along different lines
- As such, specific services may differ, but common classes exist
 - Those that provide convenience for user/programmer
 - Those that ensure efficient operations of the system

Operating System Services

- Services that provides convenience to user/programmer
 - User interface
 - Batch interface, command line interface, GUI
 - Program execution
 - Able to load, run and terminate a program
 - Input/Output operations
 - File, CD/DVD drive, Display device, SmartCard
 - File-system manipulation
 - Create/Read/Write/Delete files and directories
 - Permission management
 - Communications
 - Local/remote inter-process communication
 - Error detection
 - Able to detect error and take proper action to ensure consistent computing

Operating System Services

- Services that ensures efficient operation of the system
 - Resource allocation
 - Resources to be allocated among multiple users fairly
 - Different algorithms for different resources
 - Accounting
 - Keep track of usage statistics eg CPU, Printer, Harddisk quota
 - Reconfigure system to improve computing services
 - Protection and security
 - Security of system from outsiders
 - Ensure access to all system resources is controlled
 - Audit trail of access

Functions of an Operating System

- Regardless of services offered, there are four major group of basic functions common to all OS
 - Device management
 - Process, thread and resource management
 - Memory management
 - File management

Basic Operating System Organization

- Based on the four functions, the modern OS implements 4 major managers, which provides abstractions to the respective resource that they manage.
 - Process, Thread and Resource Manager
 - Memory Manager
 - Device Manager
 - File Manager
- Close interactions between the four managers

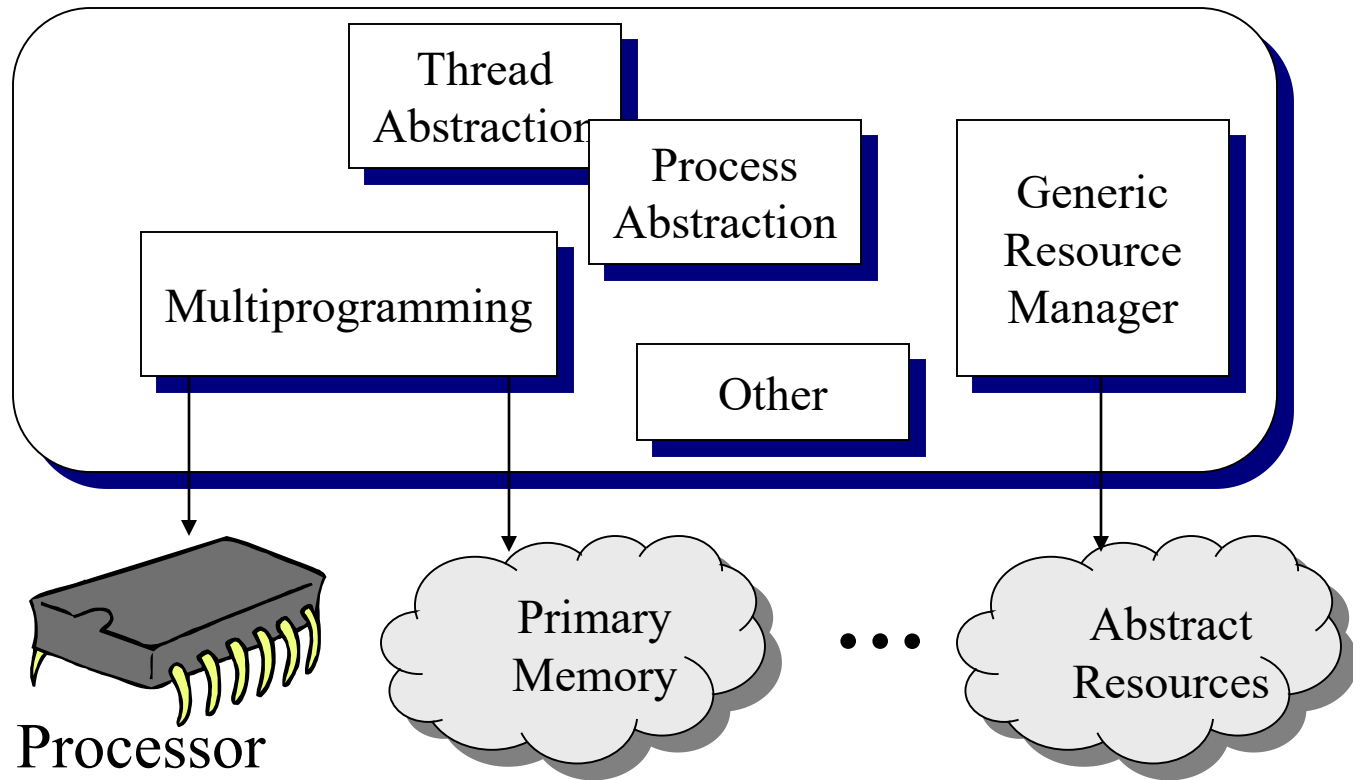
Device Management

- Refers to the way generic devices are handled.
Includes disk, tapes, terminals, printers etc
- Special management approaches for processor and memory
- Partitioning design simplifies adding and upgrading of devices

Process, Thread and Resource Management

- Creates abstractions of processes, threads, resources
- Allocates processor resource equitably
- Allocates and tracks abstract resource such as queues, semaphores, messages
- Cooperates with memory manager to administer the primary memory


Process, Thread, and Resource Management



Memory Management

- Administer and allocate primary memory
- Enforces resource isolation
- Enables sharing between processes
- Provides virtual memory extensions
 - Abstract machine's memory appear larger than physical memory

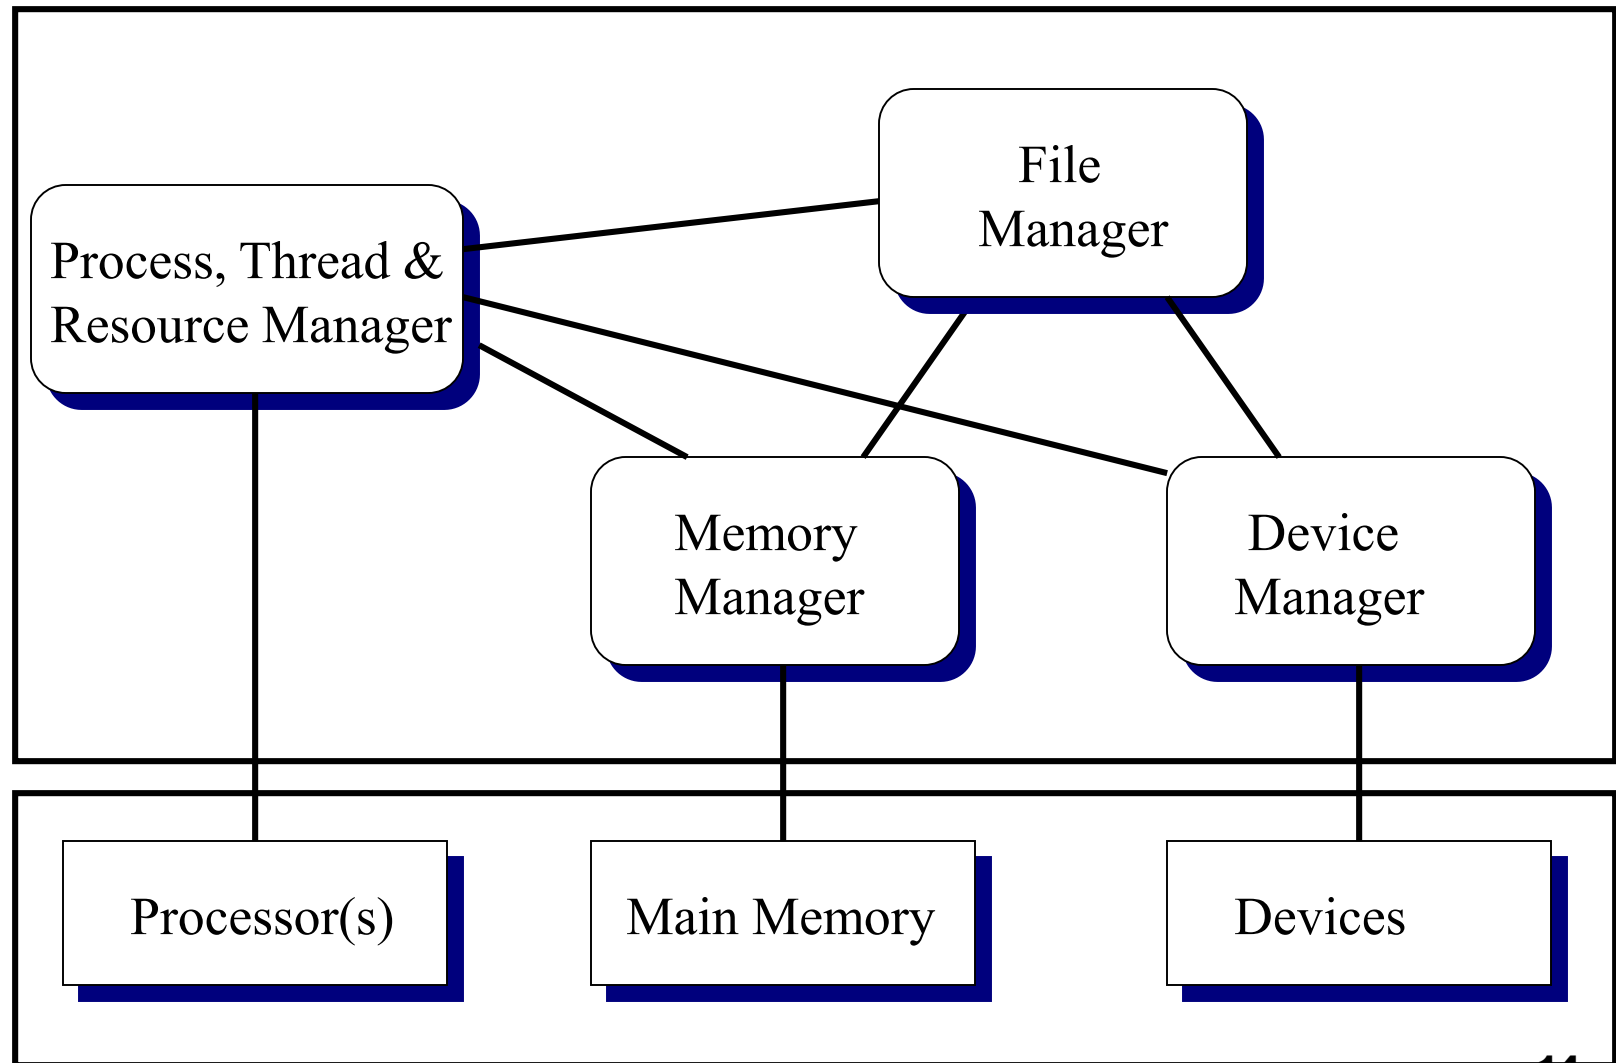
Memory Management



File Management

- Creates abstraction of storage devices i.e. I/O operations
- Range from byte stream files to indexed records
- Local and Remote file systems

Basic Operating System Organization



Key OS Requirements




Processes



*Coordinate Use
of the Abstractions*



Create the Abstractions



OS Requirements

■ Manage resource sharing

☐ Time/space-multiplexing

- Using time/space-multiplexing where appropriate.

☐ Exclusive use of a resource

- Allow processes to use a resource exclusively as required.

☐ Isolation

- Allow a resource to save information without fear of it being modified or tampered with.

☐ Managed sharing

- Sharing must be done in an orderly fashion according to the properties of the resource. E.g., printer vs disk drive.

Implementation mechanisms

- Three basic mechanisms to address isolation and sharing :
 - Processor modes
 - Kernels
 - Method of invoking system service

Processor Modes

- Distinguish between **trusted** and **un-trusted** software
- Determine execution capability and accessible memory areas
- Modern processors provide 2 modes
- Mode bit: Supervisor or User mode
- Supervisor mode (for OS)
 - Can execute all machine instructions
 - Can reference all memory locations
- User mode (for user programs)
 - Can only execute a subset of instructions
 - Can only reference a subset of memory locations

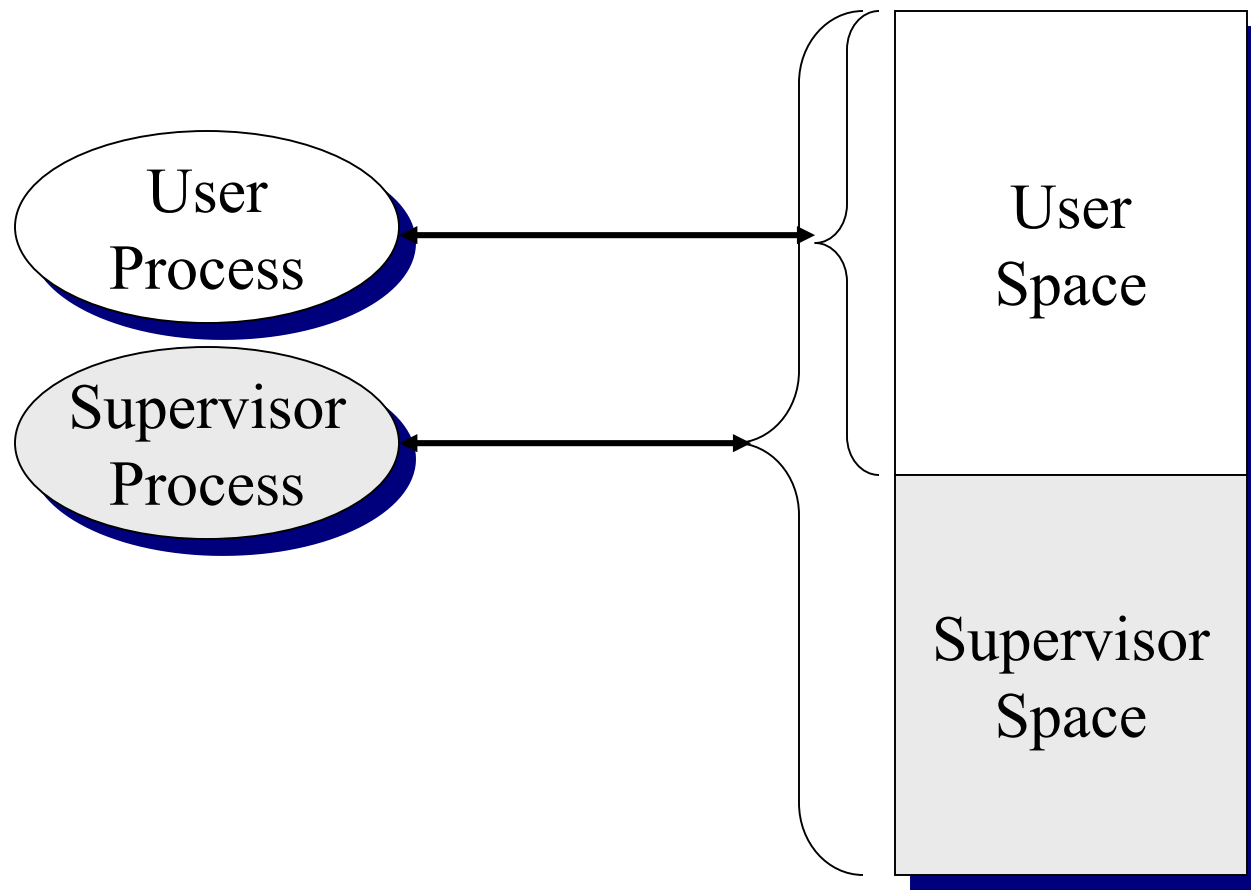
Processor Modes - Execution capability

- Processor in supervisor mode can execute all instructions
- Including **privileged** instructions. (Also called **supervisor** or **protected**.)
- Examples of privileged instructions :
 - I/O instructions
 - Memory-related instructions
 - Processor mode-change instructions
- Processor in user mode can execute only non-privileged instructions.

Processor Modes - Accessible memory

- Processor in supervisor mode can access all memory locations:
 - **System** (or **supervisor**, **kernel** or **protected**) space which refers memory area used by the OS
 - **User space** refers to memory area used by application processes
- Processor in user mode can access only user space

Processor Modes: Supervisor and User Memory



Processor Modes

- Trusted OS software executes in supervisor mode
- All other software (including some parts of the OS) executes in user mode
- Concept allows for OS to be able to control access to resources
 - User programs have to ask the OS to execute privileged instructions on their behalf.
 - Any particular configuration can isolate or permit sharing of resources according to the administrator's policy

Kernels

- The part of the OS critical to correct operation (trusted software)
- Implements the basic mechanisms that assure secure operation of entire OS
- Executes in supervisor mode
- The `trap` instruction is used to switch from user to supervisor mode, entering the OS

Requesting Services from OS

- In order to execute privileged instructions, user programs have to activate routines in the kernel, which can then execute on the user programs' behalf.
- Two techniques :
 - System call
 - Message passing

System Call

- In system call, the relevant function is activated via a **trap** instruction.
- OS provides a sub function which the user program calls
- Stub function will switch the processor to supervisor mode
- It will execute the `trap` instruction by branching to a trap table to the entry point of the system function to be invoked
- On completion, processor is switched back to user mode and control returns to user process
- Appears as ordinary function call to the application programmer


Message Passing

- In the message passing method, the user program constructs a message that requests the desired service.
- Uses OS `send()` system call
- OS kernel implements target function
- Kernel process must be started or active i.e. must be in supervisor mode, to receive message
- User process waits for result with `receive()` operation
- Kernel sends message back to user process on completion

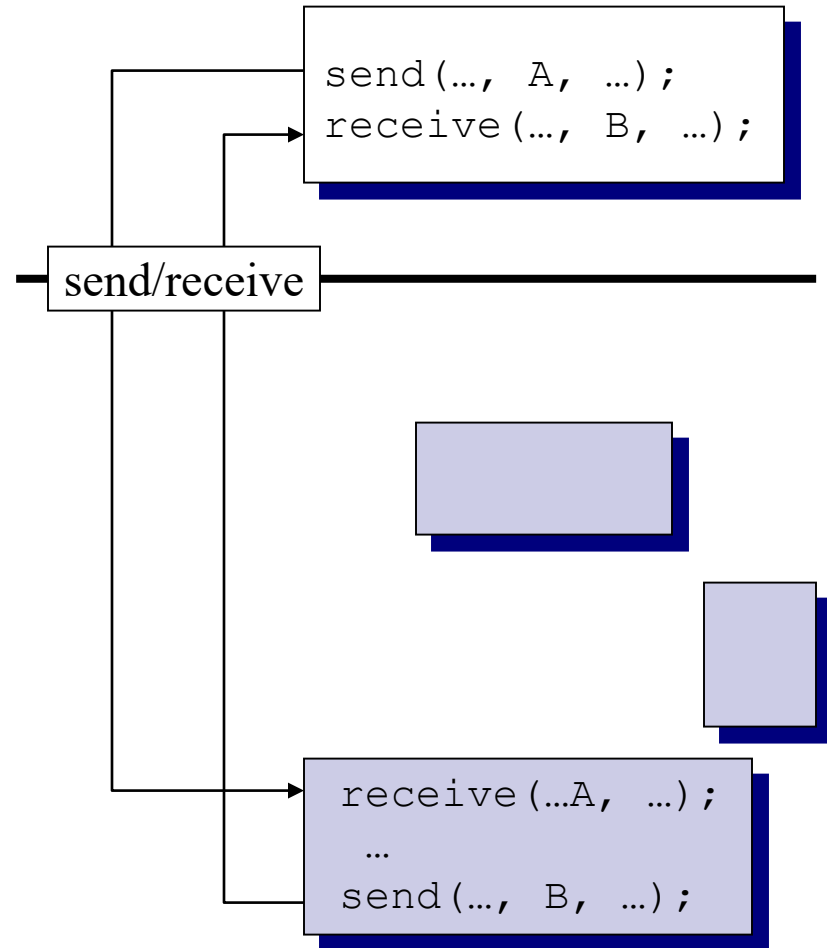
System Call versus Message Passing

- In system call, the user process/thread gains ability to execute privileged instructions
- In the message passing method, the system function is executed by the kernel process/thread
- System calls are more efficient than message passing.
 - Message passing has cost of message formation/copying and process multiplexing
 - System calls just requires a trap command.
- Most modern systems use system calls.

System Call and Message Passing Operating Systems



Procedure Call



Message Passing

System Call Using the trap Instruction

OS provides a procedure stub for user programs to execute, which will then be *translated* into the actual function via the trap table.

```
...
fork();
...

fork() {
...
trap    N_SYS_FORK()
...
}
```

Trap Table


sys_fork()

Kernel

The trap function relies on the trap table to tell it the entry point of the relevant OS function. In this case, the trap table is consulted and the address of `sys_fork()` is found so that the function can be executed.

```
sys_fork() {
/* system function */
...
return;
}
```

A Thread Performing a System Call

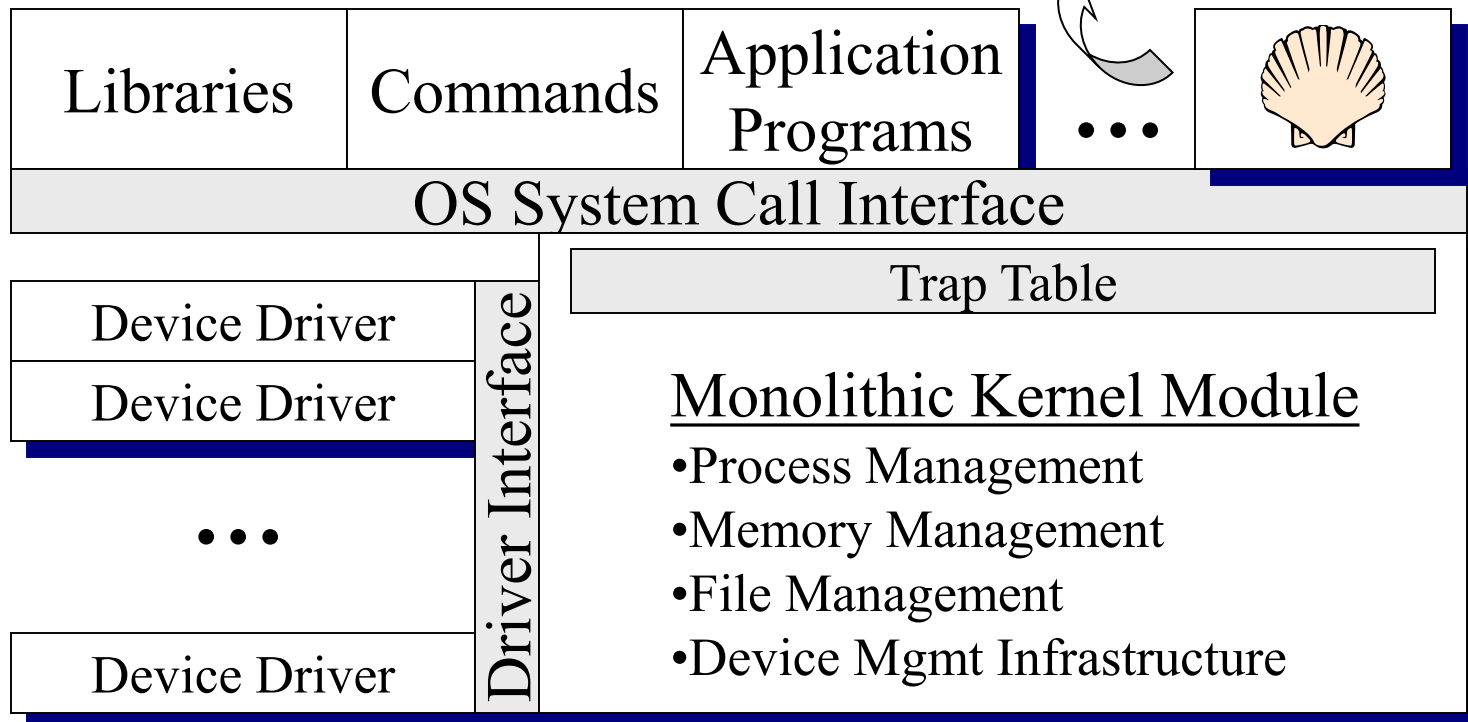



Modern OS architecture

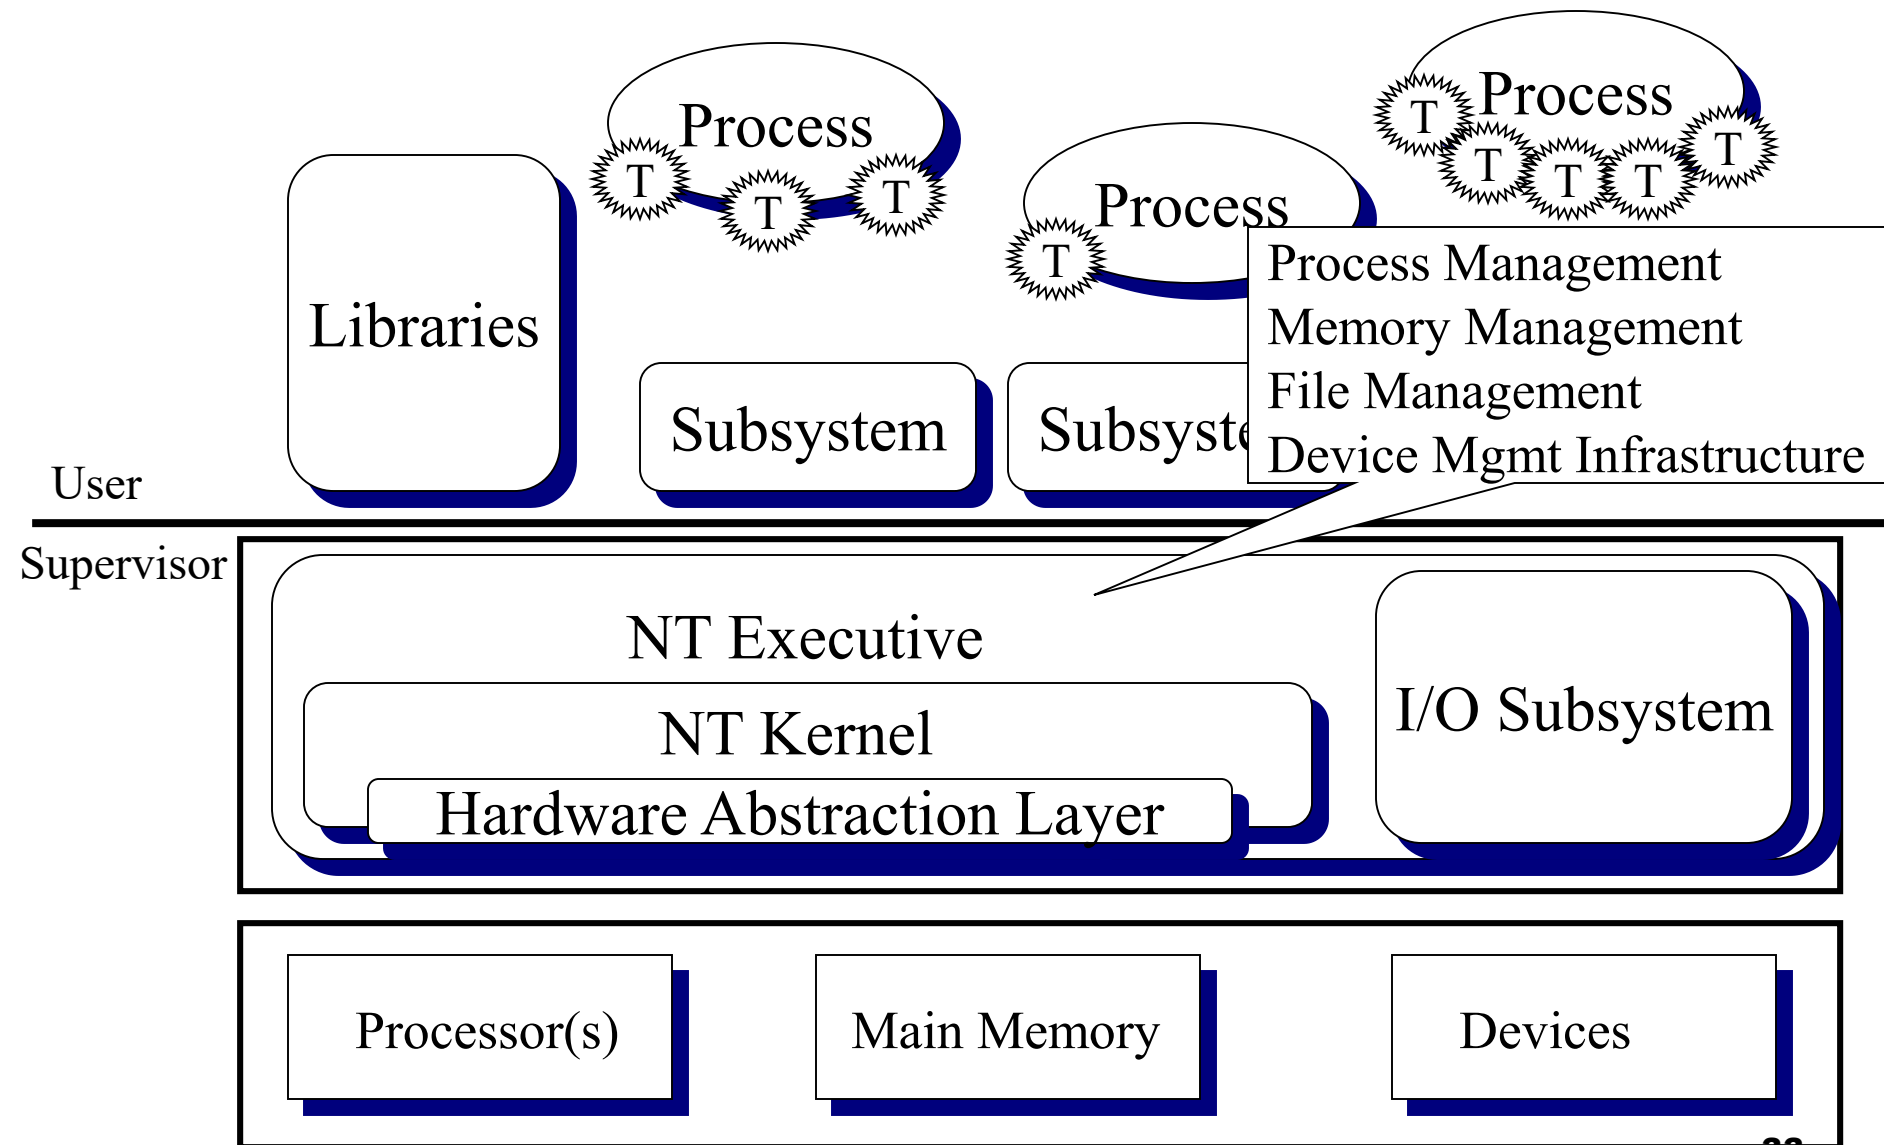
- A modular OS architecture implements each manager in its own software module
- Interaction among various managers via abstract data type
- Frequent calls incurs performance penalty
- Sacrifice modularity for performance
 - Monolithic kernel implementation
 - Four basic modules are combined into a single software module
 - Microkernel approach
 - Employs a small kernel that implements only the essential and critical functions
 - Remainder functions are implemented outside kernel, possibly in separate modules

The UNIX Architecture

Interactive User



Windows NT Organization



Conclusion

- OS implements an environment which includes processes, resources and facilities to manage resources.
- Modern OS incorporate managers for processes and resources, including memory, files and devices.
- Processor modes, kernels and a method of invoking system service allows the OS to achieve resource abstraction and sharing.