

# IT2164/IT2561

## Operating Systems and Administration

### Chapter 4


### Device Management

# Device Management


At the end of this chapter, you should be able to

- Understand device management abstraction and organization
- Understand the various I/O strategies and buffering techniques
- Understand the various device class characteristics
- To understand some of the rotating disk optimization techniques

# Input/Output Devices



Input Device



Processor



Output Device

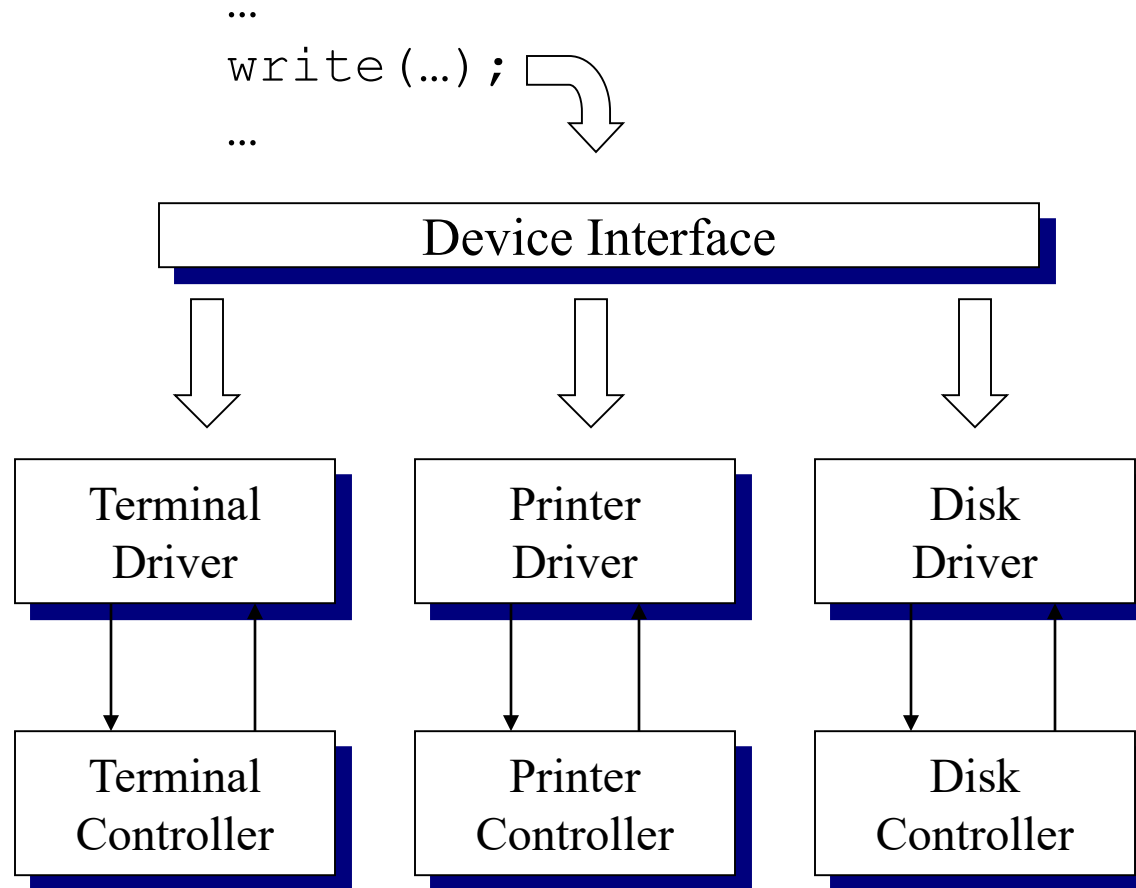
# I/O System

- In the early days of computing, input and output is at a speed similar to that of processing.
- When processors became electronic, I/O devices still remain mechanical. (e.g., hard disk vs CPU)
- Today, I/O systems are designed to handle I/O devices which are of many magnitudes slower than the slowest CPU.
  - To provide simple, abstract software interfaces to manage the I/O operations needed
  - Ensure that there is as much overlap as possible between the operation of the I/O devices and the CPU

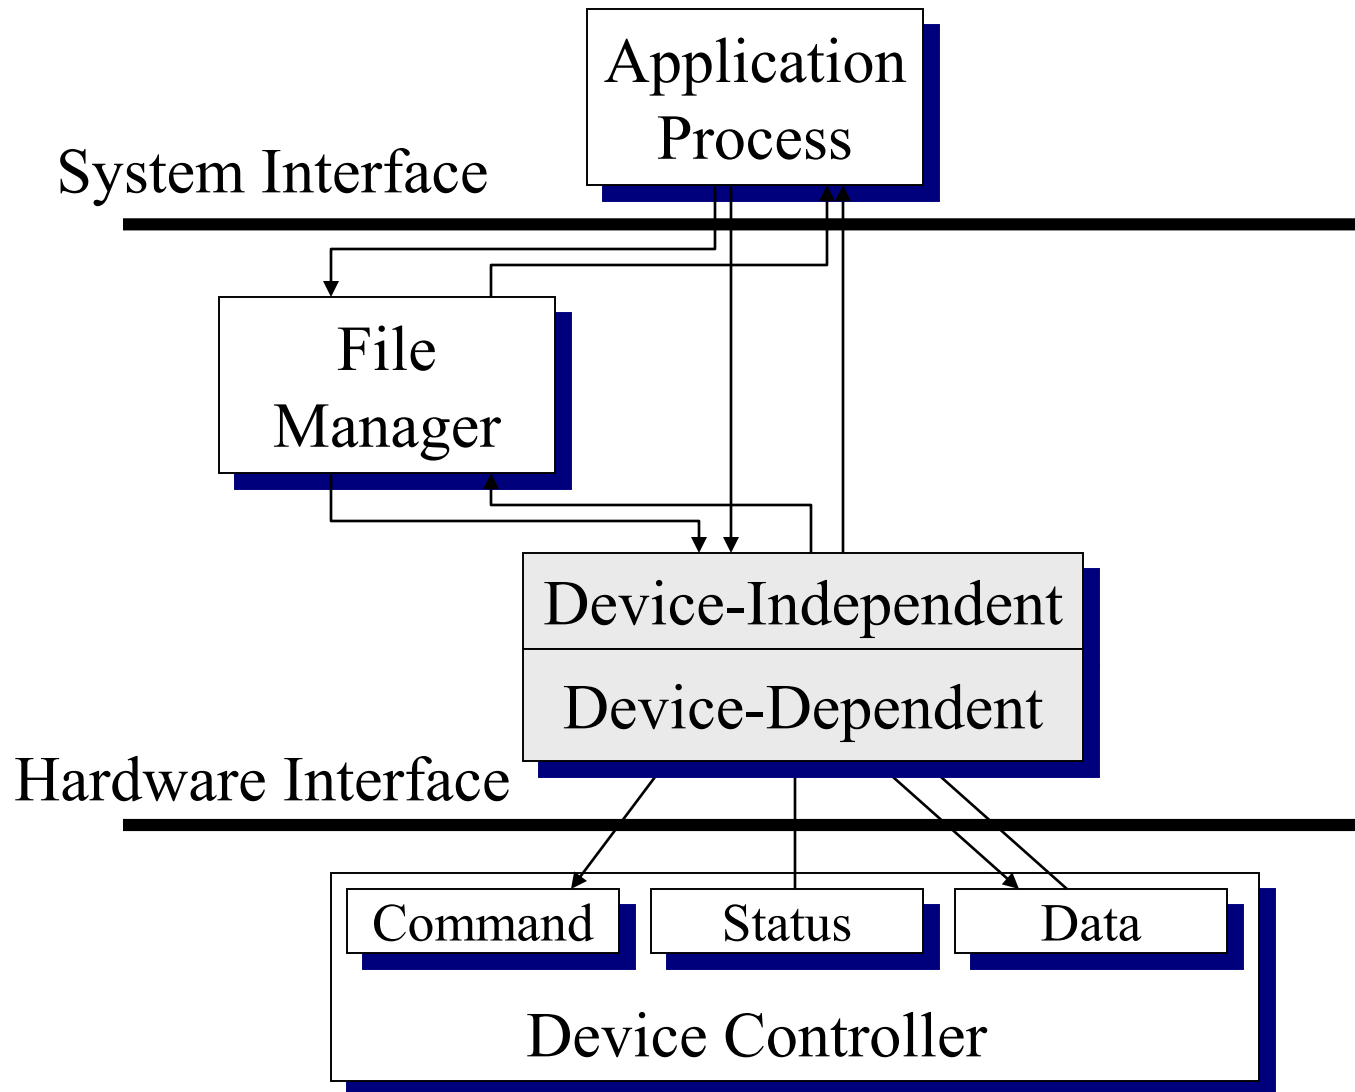
# Device Manager Abstraction

- Different devices require different operations in order to work them.
- It is important to create device drivers, which takes on the task of working these devices, but requiring only a standard set of functions to make them (the devices) work.
- The device manager, in turn, manages the collection of device drivers.
- The manager makes it possible for the OS to then provide a standard set of system calls to application programs, which use the devices.

# The Device Driver Interface



# Device Management Organization



# System Call Interface

- Functions available to application programs
- Abstract all devices (and files) to a few interfaces
- Make interfaces as similar as possible
  - Block vs character
  - Sequential vs direct access
- Device driver implements functions (one entry point per API function)



# Example: BSD UNIX Driver

open	Prepare dev for operation
close	No longer using the device
ioctl	Character dev specific info
read	Character dev input op
write	Character dev output op
strategy	Block dev input/output ops
select	Character dev check for data
stop	Discontinue a stream output op

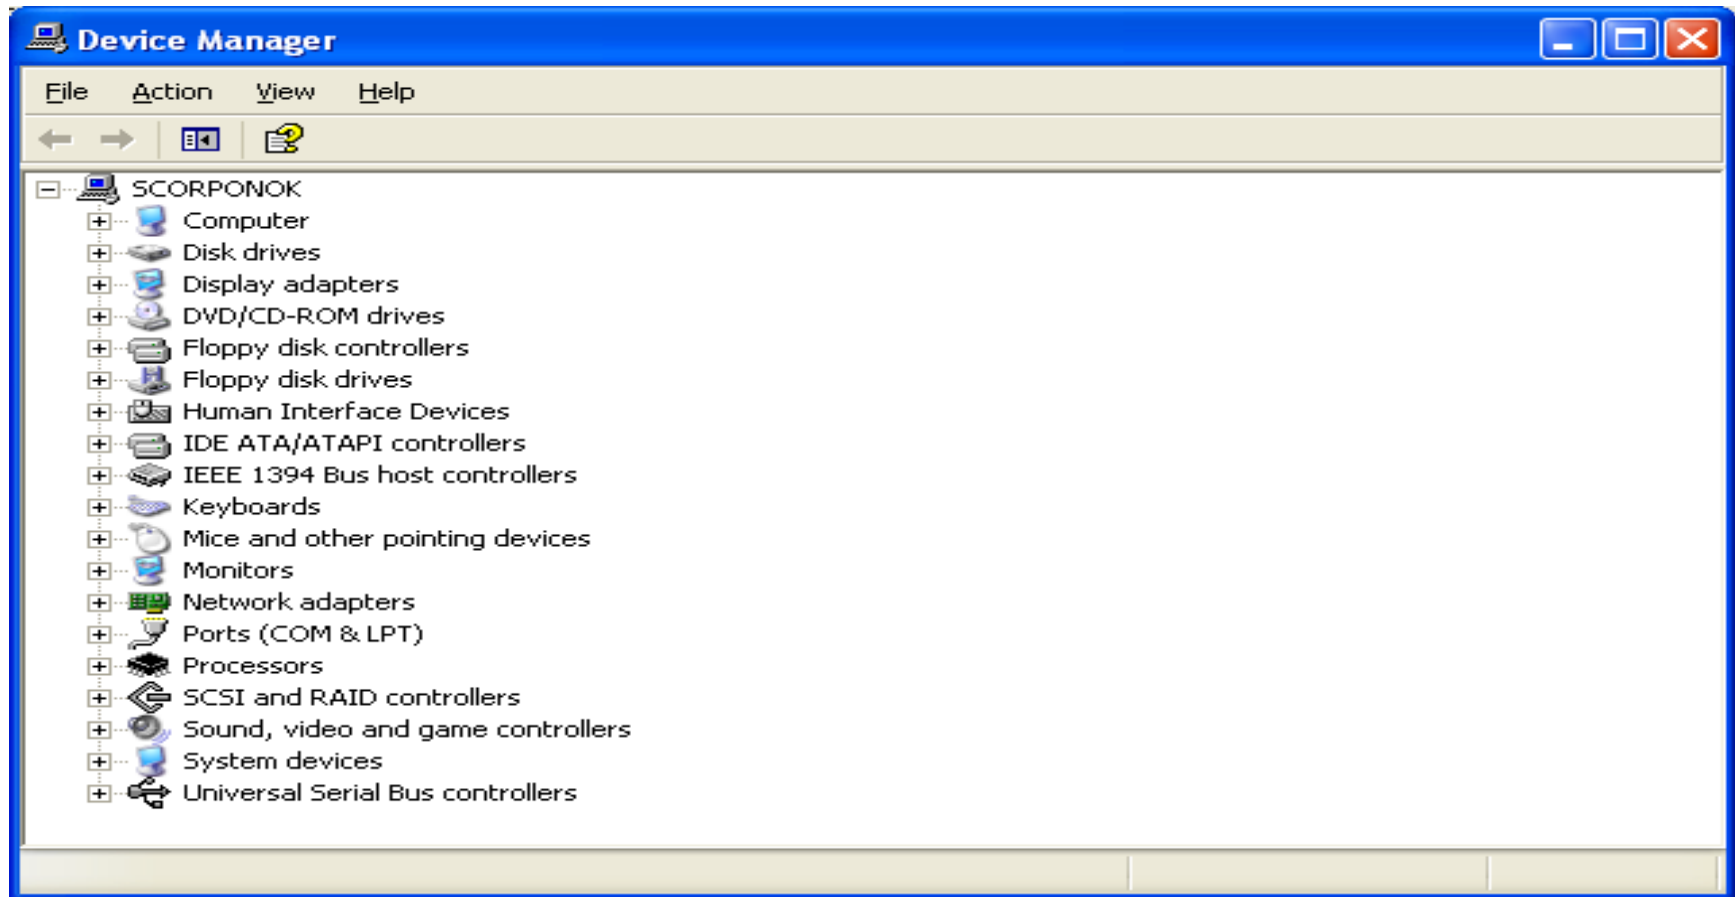
# Example: Windows system calls and low-level I/O

<code>_close</code>	Close file
<code>_commit</code>	Flush file to disk
<code>_creat, _wcreat</code>	Create file
<code>_dup</code>	Return next available file descriptor for given file
<code>_dup2</code>	Create second descriptor for given file
<code>_eof</code>	Test for end of file
<code>_lseek, _lseeki64</code>	Reposition file pointer to given location
<code>_open, _wopen</code>	Open file
<code>_read</code>	Read data from file
<code>_sopen, _wsopen, _sopen_s, _wsopen_s</code>	Open file for file sharing
<code>_tell, _telli64</code>	Get current file-pointer position
<code>_umask, _umask_s</code>	Set file-permission mask

# Device Status Table

- Device Manager keeps track of the status of the various devices using a **Device Status Table**, which consists of these info:
  - Device ID
  - Device Status (busy, done, idle)
  - Queue of processes waiting for the Device


# Device Manager in Windows



# Overlapping the Operation of a Device and the CPU

- Programs which perform IO expects IO operations to complete before the next statement is executed.
- However, performance gains can be gained if we can make the program execute instructions while the IO is taking place.
- This should be done without violating the serial execution order of the program.

# Overlapping Processing and I/O



- To maximize IO, we can have the CPU operate on another process when the IO is busy with the original process.
- This increases the efficiency of the computer and reduces the overall time required to execute all the processes.

# Buffering

- The speed of I/O is much slower than that of CPU.
- In order to speed up I/O, it is possible to keep I/O devices busy when the processes do not require I/O operations.
- This increases the overlap between I/O and processing.
- A buffer is a temporary memory-based storage area, that stores the data from an I/O operation.

# Types of Buffering


- Input buffering
  - Copy the data into memory before the process requests it.
- Output buffering
  - Temporarily stores the data in memory and have it written out to the device when the process resumes execution.
- Hardware buffering
  - Implement buffering in the hardware by having specialized registers to act as buffers. Also known as cache.
- Double buffering
  - Implement hardware buffering and having a separate buffer implemented at the software level, i.e., using the primary memory as a buffer.
- Circular buffering
  - Having a buffer which contains n number of locations. The increased number of locations allow for more data to be stored in the buffer.
  - Circular buffering is the technique of managing these locations so that they can be re-used.



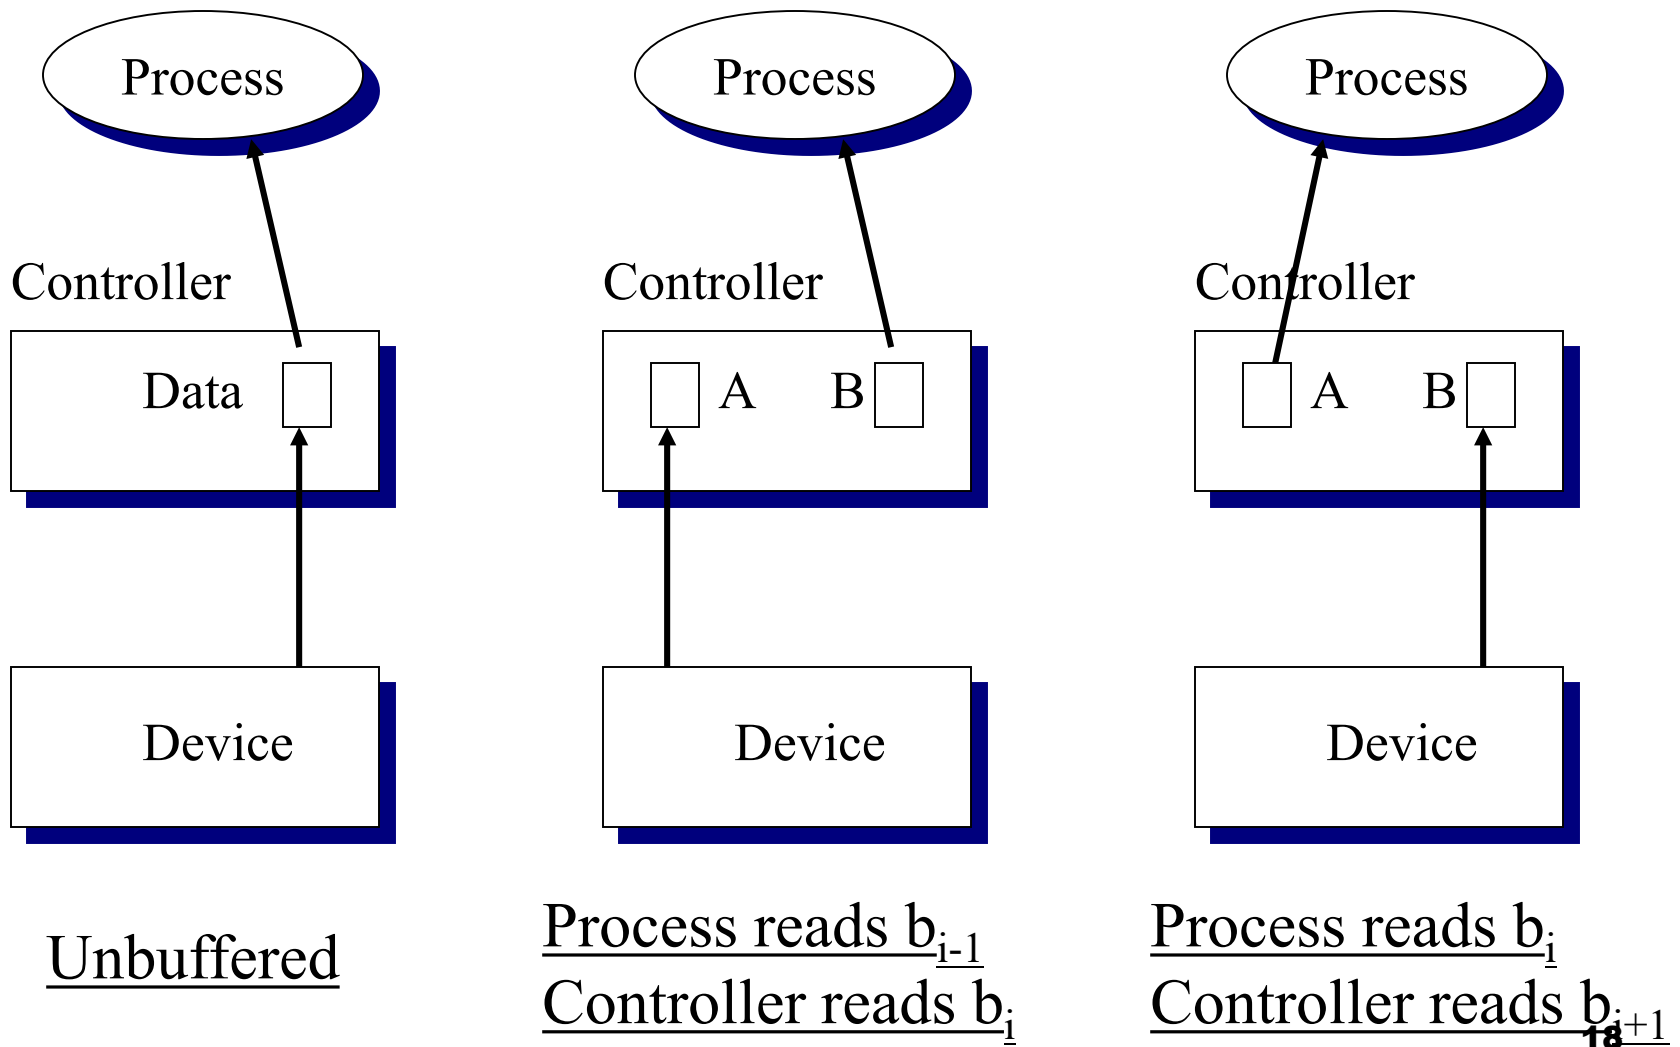
# Hardware buffer - cache

**Seagate Barracuda 7200.10 500 GB Comparison Table**

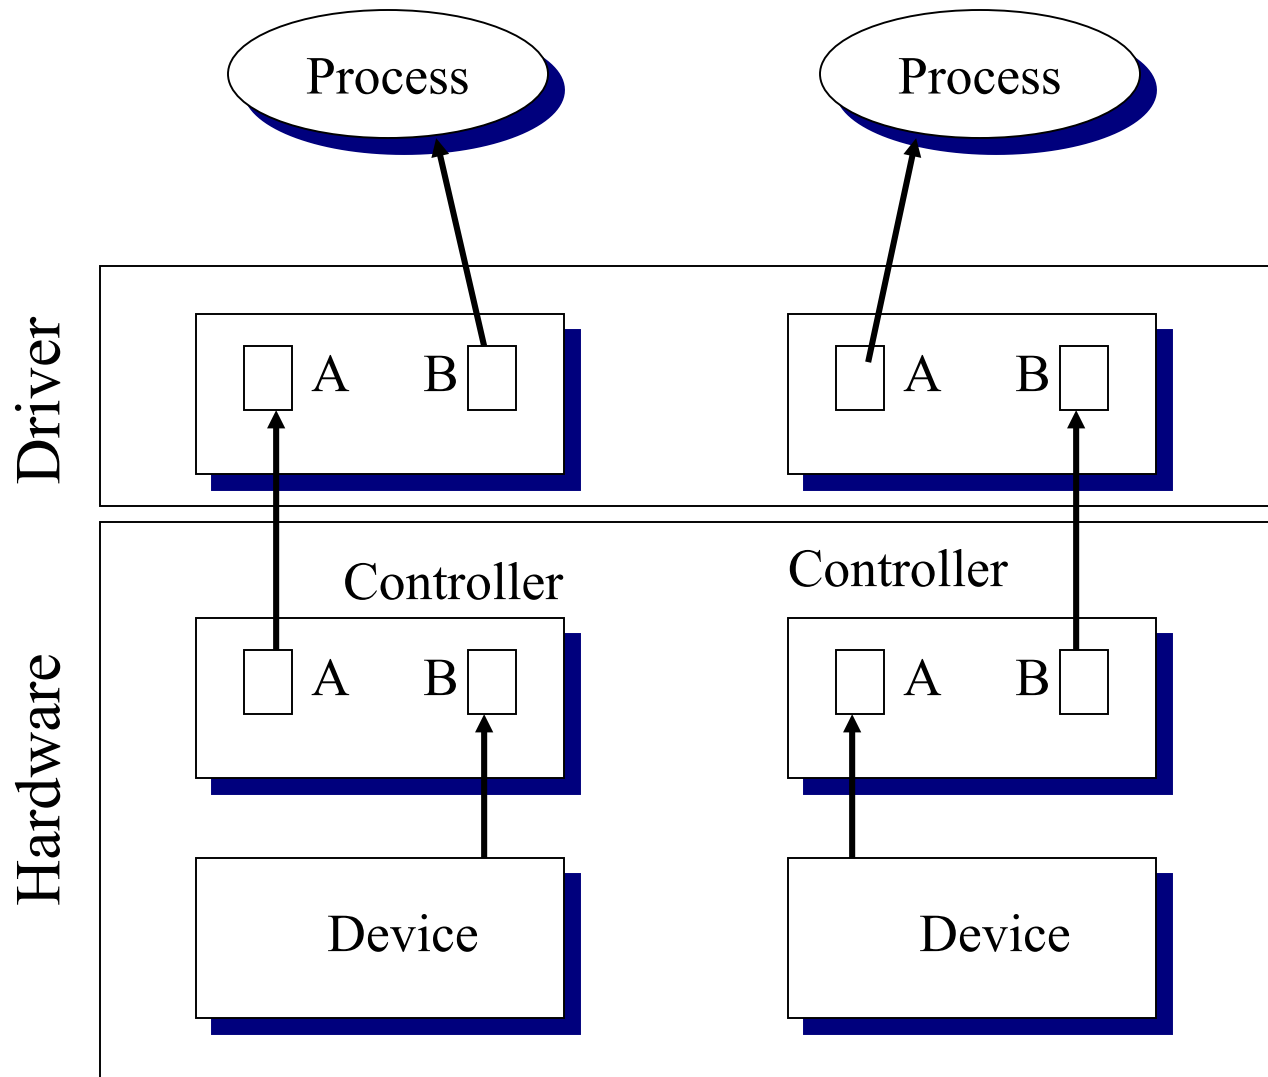
Manufacturer	Seagate	Seagate	Seagate	Seagate
Product	Barracuda	Barracuda	Barracuda	Barracuda
Model	7200.10	7200.10	7200.10	7200.10
Capacity	ST3500630A	ST3500630AS	ST3500830A	ST3500830AS
Spindle Speed	500 GB	500 GB	500 GB	500 GB
Platter	7,200 RPM	7,200 RPM	7,200 RPM	7,200 RPM
Cache	3	3	3	3
Native Command Queuing (NCQ)	16 MB	16 MB	8 MB	8 MB
Interface	No	yes	no	yes
	UltraATA/100	SATA/300	UltraATA/100	SATA/300



# Hardware Buffering



# Double Buffering in the Driver



# Device Class Characteristics

- A typical computer system will handle many types of devices, such as
  - Character devices (keyboard, monitor)
  - Block devices (printer, USB)
  - Network devices (network, modems)



- Character based devices such as tty (teletype, or terminal) and serial devices are where data stream is transferred and handled one character or byte at a time.
- Block type devices such as hard drives transfer data in blocks, typically a multiple of 256 bytes.

# Linux devices

brw-rw----	1 root disk	8,	0 Nov	7 07:06	sda
brw-rw----	1 root disk	8,	1 Nov	7 07:06	sda1
brw-rw----	1 root disk	8,	16 Nov	7 07:06	sdb
brw-rw----	1 root disk	8,	17 Nov	7 07:06	sdb1
brw-rw----	1 root disk	8,	18 Nov	7 07:06	sdb2
crw--w----	1 root tty	4,	0 Nov	7 07:06	tty0
crw--w----	1 root tty	4,	1 Nov	7 07:07	tty1
crw--w----	1 root tty	4,	10 Nov	7 07:06	tty10
crw--w----	1 root tty	4,	11 Nov	7 07:06	tty11

- Notice the leftmost character of each line in the output. The ones that have a "b" are block type devices and the ones that begin with "c" are character devices.
- sda, sda1 are naming conventions for disk devices
- tty0, tty1 are naming conventions for terminal devices or virtual consoles

# Storage Devices


- Randomly accessed storage
  - Flash memory
  - Optical disks
  - Rotating magnetic disk
  - Floppy disk
- Sequentially accessed storage
  - Magnetic tapes

# Randomly Accessed Storage Devices


- Allow a driver to access blocks of data in the device in any order. Need not be sequential.
- Non-volatile memory (such as USB memory devices) also fall in this category.
- Rotating disks are still common used today.
- Both can be easily managed using a common API.




# Rotating Disk storage



(a) Multi-surface Disk



(b) Disk Surface



(b) Cylinders

# Rotating Disk Optimizations

- Disk-based devices form the most common randomly accessed storage devices.
- E.g., Hard disks, CD/DVD, Zip disks, MO disks.
- Although they are random, we can improve on their efficiency using disk optimization techniques.
- Their effectiveness is increased in a multi-programming environment, where data blocks can be read from a large range of locations in the disk.

# Disk Optimizations

- Transfer Time: Time to copy bits from disk surface to memory
- Disk latency time: Rotational delay waiting for proper sector to rotate under R/W head
- Disk seek time: Delay while R/W head moves to the destination track/cylinder
- Access Time = seek + latency + transfer

# Disk Optimizations

- Disk optimization techniques for disk based devices
  - First-Come-First-Serve (FCFS)
  - Shortest Seek Time First (SSTF)
  - SCAN / C-SCAN
  - LOOK / C-LOOK

# Conclusion

- OS has to handle many types of devices.
- Handling alone is not enough.
- We need to find ways to optimize the use of these devices.
- Method of optimization depends on device characteristics.
- Common techniques include buffering and disk optimization (for disk-based devices).